

CMPUT 328 Fall 2024

Assignment 1

Worth 8% of the total weight

Part 1: Logistic Regression for MNIST [worth 2% of the total weight]:

Implement Logistic Regression in PyTorch (You can make use of the multiple linear regression notebook from lecture):

- Train and test on MNIST by defining your own data pipeline for training, validation and testing using PyTorch dataloader.
 - Use the last 12,000 samples of the train set as validation set.
 - Test the trained model on the validation set every few epochs to prevent overfitting.
 - Do not use the test set for training.**
 - Use Stochastic Gradient Descent ([SGD](#)) as an optimizer.
 - Use the CrossEntropy loss.
- Add a **regularization term** to improve your model (L1 or L2 regularization, whichever gives better accuracy)

Expected Performance: A correctly implemented, and somewhat well-tuned version of this algorithm will have an accuracy of **92-94%** on both test and validation sets of MNIST.

You need to complete the function *logistic_regression* in *AI_submission.py* for this part.

Part 2: Fully-connected Neural Network for CIFAR10[Worth 4% of the total weight]

Implement a Fully-connected Neural Network using the built-in functions of PyTorch. Train this network on the CIFAR10 dataset with CrossEntropy loss. The CIFAR10 dataset each image dimension is 32x32x3 as it is a color RGB image. Please refer to the *Network Architecture* section for the specifications.

You need to complete the class *FNN* in *AI_submission.py* to implement the forward pass as well as the loss computation.

- The dataset split: Training = 40000 and Validation set = 10000
- Adjust the normalization for CIFAR10
- `__init__(self, loss type, num classes)` initializes your network layers
- `forward(self, x)` takes a batch of images as a tensor of size $N \times (32 \times 32 \times 3)$ and returns the class probabilities as a tensor of size $N \times 10$ where N is the batch size
- `get_loss(self, output, target)` takes the output of the forward pass and ground truth labels of the corresponding images and returns a tensor containing the loss computed according to the loss type argument of `__init__`

Network Architecture

$$Y_p = \text{Softmax}(\text{Relu}(\text{Tanh}(XW_1 + b_1)W_2 + b_2))W_3 + b_3)$$

You can use built-in torch functions for defining the layers (`nn.Linear`) and activations. Dimensions of the vectors and matrices are as follows: X contains the input images having a shape $(N \times (32 \times 32 \times 3))$. N is the batch size. W_1 is $(32 \times 32 \times 3 \times 64)$, b_1 is (1×64) , W_2 is (64×32) , b_2 is (1×32) , W_3 is (32×10) , b_3 is (1×10) . Output probabilities Y_p has the shape $(N \times 10)$. Note that for each of N indices in the first dimension, the softmax function is applied along the second dimension of its input matrix.

Part 3: Hyperparameter Search [Worth 2% of the total weight]

Find optimal hyperparameters using Adaptive Moment Estimation ([Adam](#)) on both part 1 (Logistic Regression) (1% of the part-3 weight) and part 2 (FNN) (1% of the part-3 weight).

- You *should* perform **grid search** or **random search** for finding the optimal hyper-parameters using accuracy on the validation set and select the best configuration.
- You can also use more advanced search strategies like evolutionary search, but you are **not** allowed to use any automatic parameter search methods like [scorch](#).
- You **cannot** use the test set during this process.

You need to complete the function [tune_hyper_parameter](#) in *AI_submission.py* for this part.

Template Code

You are provided with template code in the form of three files: *AI_main.py*, *FNN_main.py* and *AI_submission.py*.

You need to complete the two functions (i.e. *logistic_regression* for part 1, *tune_hyper_parameter* for part 3) and *FNN* class in

AI_submission.py. You can add any other functions or classes you want to *AI_submission.py* but do not make any changes to *FNN_main.py* and *AI_main.py*.

Running the code

Your own machine

Install python (version ≥ 3.6) if needed and install the required packages by running:

[python3 -m pip install numpy torch torchvision tqdm paramparse](#)

Run the code using:

```
python3 AI_main.py
```

```
python3 FNN_main.py
```

It is recommended to use an IDE like pycharm or vscode to make debugging easier.

Colab

Run this from a code cell in notebooks:

```
!python3 "<full path to AI_main.py >"
```

```
!python3 "<full path to FNN_main.py >"
```

You can optionally install the *paramparse* package to enable command line arguments:

```
!python3 -m pip install paramparse
```

You can then use command line arguments as:

part 1:

```
!python3 "<full path to AI_main.py >" mode=logistic
```

part 2:

```
!python3 "<full path to FNN_main.py >" mode=fnn
```

part 3:

```
!python3 "<full path to FNN_main.py >" mode=tune target_metric=accuracy
```

Submission

You need to submit only the completed *AI_submission.py*. Make sure to import any additional libraries you need so it can be used as a standalone Python module from *FNN_main.py* and *AI_main.py*.

To reiterate, please **do not** submit *FNN_main.py* and *AI_main.py* or any other files generated by running the code.

Marking

Part 1: Marks will depend on correctness of the implementation along with the following metrics:

- **Runtime:** The total runtime of your submission (including training and testing) **should not exceed 300 seconds** for either dataset on Colab GPU.
 - One trick to improve your run time is to grid search the hyperparameters first but only put in the best hyperparameters you found in your submission.
- **Accuracy:** Score scales linearly from **83 - 93%** accuracy on the test set

Part 2: Marks will depend on correctness of the implementation along with the following metrics:

- **Runtime:** The total runtime of your submission (including training and testing) **should not exceed 300 seconds** for either dataset on Colab GPU.
- **Accuracy:** Score scales linearly from **36-40%** accuracy on the test set

Part 3: Marks will depend on the correctness of your search implementation.

- **Runtime:** The runtime of your submission **should not exceed 1500 seconds** on Colab GPU.
- **Accuracy:** There are no specific accuracy requirements except there should be improvement in loss / accuracy compared to the baseline

